

# Applying OpenMP to eigenvalue computation

Emil Løvbak and Sahar Chehrazad

December 7th, 2020

# Computing eigenvalues of a matrix $A$

- ▶ **Power method:**  $x_k = Ax_{k-1}$   
 $x_k$  converges to an eigenvector which is multiplied with the corresponding eigenvalue in each iteration.
- ▶ **Subspace iteration:**  $Q_k = AQ_{k-1}$   
We compute multiple eigenvalues by iterating with  $Q_k$ , containing columns of linearly independent vectors. We ensure independent columns of  $Q_k$  by performing a QR-decomposition with Gram-Schmidt after each iteration.

```
Q0 ← random_matrix()  
for k = 1 ... k_max do  
   $\tilde{Q}_k \leftarrow A Q_{k-1}$   
   $[Q_k, R_k] \leftarrow \text{Gram-Schmidt}(\tilde{Q}_k)$   
end for  
eigenvalues ← diagonal( $R_{k\_max}$ )
```

# Getting started

- ▶ Download the file `eigenvalues.cpp` from Toledo
- ▶ (On the VSC:) Load the module Boost
- ▶ Compile the program with `g++ -DNDEBUG -O3 -o eigenvalues eigenvalues.cpp`
- ▶ Run the program as `./eigenvalues <sizeOfA> <numberEigenvalues> <numberIterations>`
- ▶ Try to compute all eigenvalues of a 10x10 matrix with 5, 10 and 20 iterations.
- ▶ Now look at the code and try to understand how it works.

# Parallelizing with OpenMP

- ▶ Identify the different loops in the program. Can they all be parallelized (efficiently)?
- ▶ What is the most computationally expensive part of the code? What does Amdahl's law tell us?
- ▶ Play around with making different loops parallel.
  - ▶ Consider two cases:
    - ▶ Computing all eigenvalues of a (reasonably) large matrix.
    - ▶ Computing 10 eigenvalues of a very large matrix  
size of  $A \geq 10000$ .
  - ▶ Verify that you get the same results as the serial code.
  - ▶ Check what gives the best improvement to the run-time.
- ▶ (extra) Play around with scheduling of the loops, do you notice a big difference?